

# Distributed algorithms

Pierre Colson

Wednesday 05 January 2022

## Contents

<b>General</b>	<b>2</b>
Fair-loss Links . . . . .	2
Stubborn Links . . . . .	2
Reliable (Perfect) Links . . . . .	2
<b>Reliable Broadcast</b>	<b>3</b>
Best-effort Broadcast (beb) . . . . .	3
Reliable Broadcast (rb) . . . . .	3
Uniform Reliable Broadcast (urb) . . . . .	4
<b>Causal Broadcast</b>	<b>4</b>
Causality . . . . .	4
Causal Broadcast . . . . .	4
Reliable Causal Broadcast (rcb) . . . . .	5
Uniform Causal Broadcast (ucb) . . . . .	5
Reliable Causal Order Broadcast (rco) . . . . .	5
<b>Total Order Broadcast (to)</b>	<b>6</b>
<b>Consensus</b>	<b>6</b>
Algorithm 1 (cons) . . . . .	7
Algorithm 2 (ucons) . . . . .	7
Algorithm 3 . . . . .	8
<b>Shared Memory</b>	<b>9</b>
Regular Register . . . . .	9
Atomic Register . . . . .	10
<b>Terminating Reliable Broadcast (trb)</b>	<b>11</b>
<b>Non-Blocking Atomic Commit (nbac)</b>	<b>12</b>
<b>Group Membership and View Synchronous Communication</b>	<b>13</b>
Group Membership (gmp) . . . . .	13
View Synchrony (vsc) . . . . .	14

---

Markdown version on [github](#)

Compiled using [pandoc](#) and [gpdf script](#)

More fiches [here](#)

## General

- The distributed system is made of a finite set of **processes** : each process models a **sequential** program
- Every pair of processes is connected by a **link** through which the processes exchange **messages**
- **Safety** is a property which states that nothing bad should happen
- **Liveness** is a property which states that something good should happen
- Two kinds of failures are mainly considered
  - **Omissions** : The process omits to send messages it is supposed to send
  - **Arbitrary** : The process sends messages it is not supposed to send
- A **correct** process is a process that does not fail (that does not crash)
- A **Failure detector** is a distributed oracle that provides processes with suspicions about crashed processes
  - It is implemented using *timing assumptions*
  - **Perfect** :
    - \* *Strong Completeness* : Eventually, every process that crashes is permanently suspected by every other correct process
    - \* *Strong Accuracy* : No process is suspected before it crashes
  - **Eventually Perfect** :
    - \* *Strong Completeness*
    - \* *Eventually Strong Accuracy* : Eventually, no correct process is ever suspected

## Fair-loss Links

- **FL1. Fair-loss** : If a message is sent infinitely often by  $p_i$  to  $p_j$  and neither  $p_i$  or  $p_j$  crashes then  $m$  is delivered infinitely often by  $p_j$
- **FL2. Finite duplication** : If a message  $m$  is sent a finite number of times by  $p_i$  to  $p_j$ ,  $m$  is delivered a finite number of times by  $p_j$
- **FL3. No creation** : No message is delivered unless it was sent

## Stubborn Links

- **SL1. Stubborn delivery** : If a process  $p_i$  sends a message  $m$  to a correct process  $p_j$ , and  $p_i$  does not crash, then  $p_j$  delivers  $m$  an infinite number of times
- **SL2. No creation** : No message is delivered unless it was sent

```
Implements: StubbornLinks (sp2p)
Uses : FairLossLinks (flp2p)
upon event <sp2pSend, dest, m> do
  while (true) do
    trigger <flp2pSend, dest, m>
upon event <flp2pDeliver, src, m> do
  trigger <sp2pDeliver, src, m>
```

## Reliable (Perfect) Links

- **PL1. Validity** : If  $p_i$  and  $p_j$  are correct
- **PL2. No duplication** : No message is delivered (to a process) more than once
- **PL3. No creation** : No message is delivered unless it was sent
- Roughly speaking, reliable links ensure that messages exchanged between correct processes are *not lost*

```
Implements: PerfectLinks (pp2p)
Uses: StubbornLinks (sp2p)
upon event <Init> do delivered = emptySet
upon event <pp2pSend, dest, m> do
  trigger <sp2Send, dest, m>
```

```

upon event <sp2pDeliver, src, m> do
  if m not in delivered then
    trigger <pp2pDeliver, src, m>
    add m to delivered

```

## Reliable Broadcast

### Best-effort Broadcast (beb)

- **Request** : <bebBroadcast, m>
- **Indication** : <bebDeliver, src, m>
- **BEB1. Validity** : If  $p_i$  and  $p_j$  are correct then every message broadcast by  $p_i$  is eventually delivered by  $p_j$
- **BEB2. No duplication** : No message is delivered more than once
- **BEB3. No creation** : No messages is delivered unless it was broadcast

**Implements**: BestEffortBroadcast (beb)

**Uses**: PerfectLinks (pp2p)

```

upon event <bebBroadcast, m> do
  forall pi in S do
    trigger <pp2pSend, pi, m>
upon event <pp2pDeliver, pi, m> do
  trigger <dedDeliver, pi, m>

```

### Reliable Broadcast (rb)

- **Request** : <rbBroadcast, m>
- **Indication** : <rbDeliver, src, m>
- **RB1** = BEB1
- **RB2** = BEB2
- **RB3** = BEB3
- **RB4. Agreement** : For any message  $m$ , if any correct process delivers  $m$ , then every correct process delivers  $m$

**Implements**: ReliableBroadcast (rb)

**Uses**:

BestEffortBroadcast (beb)  
PerfectFailureDetector (P)

```

upon event <Init> do
  delivered = emptySet
  correct = S
  forall pi in S do from[pi] = emptySet
upon event <rbBroadcast, m> do
  delivered = delivered U {m}
  trigger <rbDeliver, self, m>
  trigger <bebBroadcast, [data, self, m]>
upon event <crash, pi> do
  correct = correct \ {pi}
  forall [pj, m] in from[pi] do
    trigger <bebBroadcast, [data, pj, m]>
upon event <bebDeliver, pi, [data, pj, m]> do
  if m not in delivered then
    delivered = delivered U {m}
    trigger <rbDeliver, pj, m>

```

```

if pi not in correct then
    trigger <bebBroadcast, [data, pj, m]>
else
    from[pi] = from[pi] U {[pj, m]}

```

## Uniform Reliable Broadcast (urb)

- **Request** : <urbBroadcast, m>
- **Indication** : <urbDeliver, src, m>
- **URB1** = BEB1
- **URB2** = BEB2
- **URB3** = BEB3
- **URB4. Uniform Agreement** : For any message  $m$ , if any process delivers  $m$ , then every correct process delivers  $m$

**Implements:** UniformBroadcast (urb)

**Uses:**

BestEffortBroadcast (beb)  
PerfectFailureDetector (P)

```

upon event <Init> do
    correct = S
    delivered = forward = emptySet
    ack[Message] = emptySet
upon event <urbBroadcast, m> do
    forward = forward U {[self, m]}
    trigger <bebBroadcast, [data, self, m]>
upon event <bebDeliver, pi, [data, pj, m]> do
    ack[m] = ack[m] U {pi}
    if [pi, m] not in forward then
        forward = forward U {[pj, m]}
        trigger <bebBroadcast, [data, pj, m]>
upon event (for any [pj, m] in forward) <correct in ack[m]> and <m not in delivered> do
    delivered = delivered U {m}
    trigger <urbDeliver, pj, m>

```

## Causal Broadcast

- A **non-blocking** algorithm using the past
- A **blocking** algorithm using **vector clocks**

### Causality

- Let  $m_1$  and  $m_2$  be any two messages :  $m_1 \rightarrow m_2$  ( $m_1$  causally precedes  $m_2$ ) iff
  - **C1. Fifo order** : Some process  $p_i$  broadcast  $m_1$  before broadcasting  $m_2$
  - **C2. Local order** : Some process  $p_i$  delivers  $m_1$  and then broadcast  $m_2$
  - **C3. Transitivity** : There is a message  $m_3$  such that  $m_1 \rightarrow m_3$  and  $m_3 \rightarrow m_2$

## Causal Broadcast

- **Request** : <coBroadcast, m>
- **Indication** : <coDeliver, src, m>
- **CO** : If any process  $p_i$  delivers a message  $m_2$ , then  $p_i$  must have delivered every message  $m_1$  such that  $m_1 \rightarrow m_2$

## Reliable Causal Broadcast (rcb)

- Request :  $\langle \text{rcoBroadcast}, m \rangle$
- Indication :  $\langle \text{rcoDeliver}, \text{src}, m \rangle$
- RB1, RB2, RB3, RB4
- CO

## Uniform Causal Broadcast (ucb)

- Request :  $\langle \text{ucoBroadcast}, m \rangle$
- Indication :  $\langle \text{ucoDeliver}, \text{src}, m \rangle$
- URB1, URB2, URB3, URB4
- CO

## Reliable Causal Order Broadcast (rco)

```
Implements: ReliableCausalOrderBroadcast (rco)
Uses : ReliableBroadcast (rb)
upon event <Init> do
    delivered = past = emptySet
upon event <rcoBroadcast, m> do
    trigger <rbBroadcast, [data, past, m]>
    past = past U {[self, m]}
upon event <rbDeliver, pi [data, pastm, m]> do
    if m not in delivered then
        forall [sn, n] in pastm do
            if n not in delivered then
                trigger <rcoDeliver, sn, n>
                delivered = delivered U {n}
                past = past U {[self, n]}
            trigger <rcoDeliver, pi, m>
            delivered = delivered U {m}
            past = past U {[pi, m]}
```

```
Implements ReliableCausalOrderBroadcast (rco)
Uses: ReliableBroadcast (rb)
upon event <Init> do
    forall pi in S: VC[pi] = 0
    pending = emptySet
upon event <rcoBroadcast, m> do
    trigger <rcoDeliver, self, m>
    trigger <rbBroadcast, [data, VC, m]>
    VC[self] = VC[self] + 1
upon event <rbDeliver, pj, [data, VCm, m]> do
    if pj not self then
        pending = pending U (pj, [data, VCm, m])
        deliver-pending
procedure deliver-pending is
    while (s, [data, VCm, m]) in pending do
        if forall pk: (VC[pk] >= VCm[pk]) do
            pending = pending - (s, [data, VCm, m])
            trigger <rcoDeliver, self, m>
            VC[s] = VC[s] + 1
```

- These algo ensure causal reliable broadcast

- If we replace reliable broadcast with uniform reliable broadcast, these algo would ensure uniform causal broadcast

## Total Order Broadcast (to)

- In **reliable** broadcast, the processes are free to deliver messages in any order they wish
- In **causal** broadcast, the processes need to deliver messages according to some order (causal order)
  - The order imposed by causal broadcast is however partial : some messages might be delivered in different order by the processes
- In **total order** broadcast, the processes must deliver all messages according to the same order (i.e. the order is now total)
  - This order does not need to respect causality (or event FIFO ordering)
- **Request** :  $\langle \text{toBroadcast}, m \rangle$
- **Indication** :  $\langle \text{toDeliver}, \text{src}, m \rangle$
- **RB1. Validity** : If  $p_i$  and  $p_j$  are correct, then every message broadcast by  $p_i$  is eventually delivered by  $p_j$
- **RB2. No duplication** : No message is delivered more than once
- **RB3. No creation** : No message is delivered unless it was broadcast
- **RB4. (Uniform) Agreement** : For any message  $m$ . If a correct (any) process delivers  $m$ , then every correct process delivers  $m$
- **(Uniform) Total order** : Let  $m$  and  $m'$  be any two messages. Let  $p_i$  be any (correct) process that delivers  $m$  without having delivered  $m'$ . Then no (correct) process delivers  $m'$  before  $m$
- Uses consensus (see next chapter)

**Implements:** TotalOrder (to)

**Uses:**

```

ReliableBroadcast (rb)
Consensus (cons)
upon event <Init> do
  unordered = delivered = emptySet
  wait = false;
  sn = 1
upon event <toBroadcast, m> do
  trigger <rbBroadcast, m>
upon event <rbDeliver, sm, m> and (m not in delivered) do
  unordered = unordered U {(sm, m)}
upon event (unordered not emptySet) and not wait do
  wait = true
  trigger <Propose, unordered>sn
upon event <Decide, decided>sn do
  unordered = unordered \ decided
  ordered = deterministicSort(decided)
  forall (sm, m) in ordered do
    trigger <toDeliver, sm, m>
    delivered = delivered U {m}
  sn = sn + 1
  wait = false

```

## Consensus

- In the (uniform) consensus problem the processes propose values and need to agree on one among these values
- **Request** :  $\langle \text{Propose}, v \rangle$

- **Indication** :  $\langle \text{Decide}, v' \rangle$
- **C1. Validity** : Any value decided is a value proposed
- **C2. (Uniform) Agreement** : No two correct (any) processes decide differently
- **C3. Termination** : Every correct process eventually decides
- **C4. Integrity** : Every process decides at most once, No process decides twice

### Algorithm 1 (cons)

- A  $P$ -based (fail-stop) consensus algorithm
- The processes exchange and update proposals in rounds and decide on the value of the non-suspected process with the smallest id
- The processes go through rounds incrementally (1 to  $n$ ) : in each round, the process with the id corresponding to that round is the leader of the round
- The leader of a round decides its current proposal and broadcast it to all
- A process that is not leader in a round waits to deliver the proposal of the leader in that round to adopt, or to suspect the leader

**Implements:** Consensus (cons)

**Uses:**

BestEffortBroadcast (beb)  
PerfectFailureDetector (P)

```

upon event <Init> do
  suspected = emptySet
  round = 1
  currentProposal = null
  broadcast = delivered[] = false
upon event <crash, pi> do
  suspected := suspected U {pi}
upon event <Propose, v> do
  if currentProposal = null then
    currentProposal = v
upon event <bebDeliver, p_round, value> do
  currentProposal = value
  delivered[round] = true
upon event delivered[round] = true or p_round in suspected do
  round = round + 1
upon event p_round = self and broadcasted = false and currentProposal not null do
  trigger <Decide, currentProposal>
  trigger <bebBroadcast, currentProposal>
  broadcast = true

```

### Algorithm 2 (ucons)

- A  $P$ -based (fail-stop) uniform consensus algorithm
- The processes exchange and update proposal in rounds, and after  $n$  rounds decide on the current proposal value
- Implements uniform consensus
- The processes go through rounds incrementally (1 to  $n$ ): in each round  $I$ , process  $p_I$  sends its current-Proposal to all
- A process adopts any currentProposal it receives
- Processes decide on their currentProposal values at the end of round  $n$

**Implements:** Uniform Consensus (ucons)

**Uses :**

BestEffortBroadcast (beb)

```

PerfectFailureDetector (P)
upon event <Init> do
    suspected = emptySet
    round = 1
    currentProposal = null
    broadcast = delivered[] = false
    decided = false
upon event <crash, pi> do
    suspected = suspected U {pi}
upon event <Propose, v> do
    if currentProposal = null then
        currentProposal = v
upon event <bebDeliver, p_round, value> do
    currentProposal = value
    delivered[round] = true
upon event delivered[round] = true or p_round in suspected do
    if round = n and decided = false then
        trigger <Decide, currentProposal>
        decided = true
    else
        round = round + 1
upon event p_round = self and broadcast = false and currentProposal not null do
    trigger <bebBroadcast, currentProposal>
    broadcast = true

```

### Algorithm 3

- A  $\langle \rangle P$ -based uniform algorithm assuming a correct majority
- The processes alternate in the role of a coordinator until one of them succeeds in imposing a decision
- A uniform consensus algorithm assuming
  - A correct majority
  - A  $\langle \rangle P$  failure detector
- The processes alternate in the role of a phase coordinator until one of them succeeds in imposing a decision
- $\langle \rangle P$  ensures
  - **Strong completeness** : eventually every process that crashes is permanently suspected by all correct processes
  - **Eventual strong accuracy** : eventually no correct process is suspected by any process
    - \* Strong accuracy holds only after finite time
  - Correct processes may be falsely suspected a finite number of times
  - This breaks consensus algorithm 1 and 2
- This algorithm is also round based : process move incrementally from one round to the other
- Process  $p_i$  is leader in every round  $k$  such that  $k \bmod N = i$
- In such a round,  $p_i$  tries to decide
  - $p_i$  succeeds if it is not suspected (process that suspect  $p_i$  inform  $p_i$  and move to the next round,  $p_i$  does so as well)
  - If  $p_i$  succeeds,  $p_i$  uses a reliable broadcast to send the decision to all (the reliability of the broadcast is important here to preclude the case where  $p_i$  crashes, some other processes delivers the message and stop while rest keeps going without majority)
- 1.  $p_i$  selects among a majority the latest adopted value (latest with respect to the round in which the value is adopted)
- 2.  $p_i$  imposes that value at a majority : any process in that majority adopts that value -  $p_i$  fails if it is suspected
- 3.  $p_i$  decides and broadcasts the decision to all



# Shared Memory

## Regular Register

- Assumes only one writer
- Provides *strong* guarantees when there is no concurrent operations
- When some operations are concurrent, the register provides *minimal* guarantees
- `Read()` returns :
  - The last value written if there is no concurrent or failed operations
  - Otherwise the last value written on *any* value concurrently written i.e. the input parameter of some `Write()`
- We assume **fail-stop** model
  - Process can fail by crashing (no recovery)
  - Channels are reliable
  - Failure detection is perfect
- We implement a **regular** register
  - Every process  $p_i$  has a local copy of the register value  $v_i$
  - Every process reads **locally**
  - The writer writes **globally**

```
Write(v) at pi
send [W, w] to all
forall pj, wait until either
    receive [ack] or
    detect [pj]
return ok
```

```
Read() at pi
return vi
```

```
At pi
when receive [W, w] from pj
    vi = v
    send [ack] to pj
```

- We assume while failure detection is not perfect
  - $P_1$  is the writer and any process can be reader
  - A majority of the process is correct
  - Channels are reliable
- We implement a **regular** register
  - Every process  $p_i$  maintains a local copy of the register  $v_i$ , as well as a sequence number  $sn_i$  and a read timestamp  $rs_i$
  - Process  $p_1$  maintains in addition a timestamp  $ts_1$

```
Write(v) at p1
ts1 ++
send [W, ts1, v] to all
when receive [W, ts1, ack] from majority
    return ok
```

```
Read() at pi
rsi ++
send [R, rsi] to all
when receive [R, rsi, snj, vj] from majority
    v = vj with the largest snj
return v
```

```

At pi
  when receive [W, ts1, v] from p1
    if ts1 > sni then
      vi = v
      sni = ts1
      send[W, ts1, ack] to p1
  when receive [R, rsj] from pj
    send [R, rsj, sni, vi] to pj

```

## Atomic Register

- An **Atomic Register** provides strong guarantees even when there is concurrency and failures : the execution is equivalent to a sequential and failure-free execution
- Every failed (write) operation appears to be either complete or not to have been invoked at all
- Every complete operation appears to be executed at some instant between its invocation and reply time events
- We implement a **fail-stop 1-N atomic register**
  - Every process maintains a local value of the register as well as a sequence number
  - The writer,  $p_1$ , maintains, in addition a timestamp  $ts_1$
  - Any process can read in the register

```

Write(v) at p1
  ts1++
  send [W, ts1, v] to all
  forall pi wait until either
    receive [ack] or
    detect [pi]
  return ok

```

```

Read() at pi
  send [W, sni, vi] to all
  forall pi wait until either
    receive [ack] or
    suspect [pj]
  return vi

```

```

At pi
  When pi receive [W, ts, v] from pj
    if ts > sni then
      vi = v
      sni = ts
      send [ack] to pj

```

- We implement a **fail-stop N-N atomic register**

```

Write(v) at pi
  send [W] to all
  forall pj wait until either
    receive [W, snj] or
    suspect [pj]
  (sn, id) = (highest snj + 1, i)
  send [W, (sni, id), v] to all
  forall pj wait until either
    receive [W, (sn, id), ack] or

```

```

    detect [pj]
return ok

```

```

Read() at pi
send [R] to all
forall pj wait until either
    recieve [R, (snj, idj), vj] or
    suspect pj
v = vj with the highest (snj, idj)
(sn, id) = highest (snj, idj)
send [W, (sn, id), v] to all
forall pj wait until either
    receive [W, (sn, id), ack] or
    detect [pj]
return v

```

```

At pi
T1 :
    when receive [W] from pj
        send [W, sn] to pj
    when receive [R] from pj
        send [R, (sn, id), vi] to pj

T2 :
    when receive [W, (snj, idj), v] from pj
        if (snj, idj) > (sn, id) then
            vi = v
            (sn, id) = (snj, idj)
        send [W, (sn, id), ack] to pj
    when receive [W, (snj, idj), v] from pj
        if (snj, idj) > (sn, id) then
            vi = v
            (sn, id) = (snj, idj)
        send [W, (sn, id), ack] to pj

```

- From fail-stop to **fail-silent**
  - We assume a majority of correct processes
  - In the 1-N algorithm, the writer writes in a majority using a timestamp determined locally and the reader selects a value from a majority and then imposes this value on a majority
  - In the N-N algorithm, the writers determines first the timestamp using a majority

## Terminating Reliable Broadcast (trb)

- Like reliable broadcast, terminating reliable broadcast (TRB) is a communication primitive used to disseminate a message among a set of processes in a reliable way
- TRB is however strictly stronger than (uniform) reliable broadcast
- Like with reliable broadcast, correct processes in TRB agree on the set of messages they deliver
- Like with (uniform) reliable broadcast, every correct process in TRB delivers every message delivered by any correct process
- Unlike with reliable broadcast, every correct process delivers a message, event if the broadcaster crashes
- The problem is defined for a specific broadcaster process  $p_i = src$  (known by all processes)
  - Process  $src$  is supposed to broadcast a message  $m$  (distinct from  $\varphi$ )
  - The other processes need to deliver  $m$  if  $src$  is correct but may deliver  $\varphi$  if  $src$  crashes
- **Request** :  $\langle trbBroadcast, m \rangle$

- **Indication** :  $\langle \text{trbDeliver}, p, m \rangle$
- **TRB1. Integrity** : If a process delivers a message  $m$ , then either  $m$  is  $\varphi$  or  $m$  was broadcasted by  $src$
- **TRB2. Validity** : If the sender  $src$  is correct and broadcasts a message  $m$ , then  $src$  eventually delivers  $m$
- **TRB3. (Uniform) Agreement** : For any message  $m$ , if a correct (any) process delivers  $m$ , then every correct process delivers  $m$
- **TRB4. Termination** : Every correct process eventually delivers exactly one message

**Implements:** `trbBroadcast (trb)`

**Uses:**

BestEffortBroadcast (beb)  
PerfectFailureDetector (P)  
Consensus (cons)

```

upon event <Init> do
  prop = 0
  correct = S
upon event <trbBroadcast, m> do
  trigger <bebBroadcast, m>
upon event <crash, src> and (prop = 0) do
  prop = phi
upon event <bebDeliver, src, m> and (prop = 0) do
  prop = m
upon event (prop not 0) do
  trigger <Propose, prop>
upon event <Decide, decision> do
  trigger <trbDeliver, src, decision>

```

- We give an algorithm that implements  $P$  unsigned TRB. More precisely, we assume that every process  $p_i$  can use an infinite number of instances of TRB where  $p_i$  is the sender  $src$ 
  1. Every process  $p_i$  keeps on trbBroadcasting messages  $m_{i1}, m_{i2}$  etc
  2. If a process  $p_k$  delivers  $\varphi_i$ ,  $p_k$  suspects  $p_i$

## Non-Blocking Atomic Commit (nbac)

- A **transaction** is an atomic program describing a sequence of accesses to shared and distributed information
  - Can be determined either by **committing** or **aborting**
- **Atomicity** : a transaction either performs entirely or none at all
- **Consistency** : a transaction transforms a consistent state into another consistent state
- **Isolation** : a transaction appears to be executed in isolation
- **Durability** : the effects of a transaction that commits are permanent
- As in consensus, every process has an initial value 0 (no) or 1 (yes) and must decide on a final value 0 (abort) or 1 (commit)
- The proposition means the ability to commit the transaction
- The decision reflects the contract with the user
- Unlike consensus, the processes here seek to decide 1 but every process has a veto right
- **Request** :  $\langle \text{Propose}, v \rangle$
- **Indication** :  $\langle \text{Decide}, v' \rangle$
- **NBAC1. Agreement** : No two processes decide differently
- **NBAC2. Termination** : Every correct process eventually decides
- **NBAC3. Commit-Validity** : 1 can only be decided if all process propose 1
- **NBAC4. Abort-Validity** : 0 can only be decided if some process crashes or votes 0

**Implements:** `nonBlockingAtomicCommit (nbac)`

**Uses:**

```

BestEffortBroadcast (beb)
PerfectFailureDetector (P)
UniformConsensus (uniCons)
upon event <Init> do
  prop = 1
  delivered = emptySet
  correct = pi
upon event <crash, pi> do
  correct = correct \ {pi}
upon event <Propose, v> do
  trigger <bebBroadcast, v>
upon event <bebDeliver, pi, v> do
  delivered = delivered U {pi}
  prop = prop * v
upon event correct \ deliver = empty do
  if correct not pi then
    prop = 0
  trigger <uncPropose, prop>
upon event <uncDecide, decision> do
  trigger <Decide, decision>

```

## Group Membership and View Synchronous Communication

### Group Membership (gmp)

- In many distributed applications, processes need to know which processes are **participating** in the computation and which are not
- Failure detector provide such information; however that information is **not coordinated** event if the failure detector is perfect
- Like with a failure detector, the processes are informed about failures, we say that the processes **install views**
- Like with a perfect failure detector, the processes have accurate knowledge about failures
- Unlike with a perfect failure detector, the information about failures are **coordinated** : the processes install the same sequence of views
- **Indication** :  $\langle \text{membView}, V \rangle$
- **Memb1. Local Monotonicity** : if a process installs view  $(j, M)$  after installing  $(k, N)$ , then  $j > k$  and  $M \leq N$
- **Memb2. Agreement** : no two processes install views  $(j, M)$  and  $(j, M')$  such that  $M \neq M'$
- **Memb3. Completeness** : if a process  $p$  crashes, then there is an integer  $j$  such that every correct process eventually installs view  $(j, M)$  such that  $p \notin M$
- **Memb4. Accuracy** : if some process installs view  $(i, M)$  and  $p \notin M$  then  $p$  has crashed

**Implements:** groupMembership (gmp)

**Uses:**

```

PerfectFailureDetector (P)
UniformConsensus (Ucons)
upon event <Init> do
  view = (0, S)
  correct = S
  wait = true
upon event <crash, pi> do
  correct = correct \ {pi}
upon event (correct < view.memb) and (wait = false) do

```

```

wait = true
trigger <ucPropose, (view.id + 1, correct)>
upon event <ucDecide, (id, memb)> do
  view = (id, memb)
  wait = false
  trigger <membView, view>

```

## View Synchrony (vsc)

- **View synchronous broadcast** is an abstraction that results from the combination of group membership and reliable broadcast
  - Ensures that the delivery of messages is coordinated with the installation of views
- **Request** : <vsBroadcast, m> (<vsBlockOk>)
- **Indication** : <vsDeliver, src, m>, <vsView, V> (<vsBlock>)
- **Memb1, Memb2, Memb3, Memb4, RB1, RB2, RB3, RB4**
- **VS** : a message is **vsDelivered** in the view where it is **vsBroadcast**
- If the application keeps **vsBroadcasting** messages, the **view synchrony** abstraction might never be able to **vsInstall** a new view, the abstraction would be impossible
  - We introduce a specific event for the abstraction to block the application from **vsBroadcasting** messages, this only happens when a process crashes

**Implements:** ViewSynchrony (vs)

**Uses:**

GroupMembership (gmp)  
 TerminationReliableBroadcast (trb)  
 BestEffortBroadcast (beb)

```

upon event <Init> do
  view = (0, S)
  nextView = ()
  pending = delivered = trbDone = emptySet
  flushing = blocked = false
upon event <vsBroadcast, m> and (blocked = false) do
  delivered = delivered U {m}
  trigger <vsDeliver, self, m>
  trigger <bebBroadcast, [data, view.id, m]>
upon event <bebDeliver, src, [data, vid, m]> do
  if (view.id = vid) and (m not in delivered) and (blocked = false) then
    delivered = delivered U {m}
    trigger <vsDeliver, src, m>
upon event <membView, V> do
  addtoTail(pending, V)
upon event (pending not emptySet) and (flushing = false) do
  nextView = removeFromHead(pending)
  flushing = true
  trigger <vsBlock>
upon event <vsBlockOk> do
  blocked = true
  trbDone = emptySet
  trigger <trbBroadcast, self, (view.id, delivered)>
upon event <trbDeliver, p, (vid, del)> do
  trbDone = trbDone U {p}
  forall m in del and m not in delivered do
    delivered = delivered U {m}
    trigger <vsDeliver, src, m>

```

```

upon event (trbDone = view.memb) and (blocked = true) do
  view = nextView
  flushing = blocked = false
  delivered = emptySet
  trigger <vsView, view>

```

Implements: ViewSynchrony (vs)

Uses:

```

  UniformConsensus (uc)
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)

```

```

upon event <Init> do
  view = (0, S)
  correct = S
  flushing = blocked = false
  delivered = dset = emptySet
upon event <vsBroadcast, m> and (blocked = false) do
  delivered = delivered U {m}
  trigger <vsDeliver, self, m>
  trigger <bebBroadcast, [data, view.id, m]>
upon event <bebDeliver, src, [data, vid, m]> do
  if (view.id = vid) and (m not in delivered) and (blocked = false) then
    delivered = delivered U {m}
    trigger <vsDeliver, src, m>
upon event <crash, p> do
  correct = correct \{p}
  if flushing = false then
    flushing = true
    trigger <vsBlock>
upon event <vsBlockOk> do
  blocked = true
  trigger <bebBroadcast, [DSET, view.id, delivered]>
upon event <bebDeliver, src, [DSET, vid, del]> do
  dset = dset U (src, del)
  if forall p in correct, (p, mset) in dset then
    trigger <ucPropose, view.id+1, correct, dset>
upon event <ucDecided, id, memb, vsdset> do
  forall (p, mset) in vsdest and p in memb do
    forall (src, m) in mset and m not in delivered do
      delivered = delivered U {m}
      trigger <vsDeliver, src, m>
  viewx = (id, memb)
  flushing = blocked = false
  dset = delivered = emptySet
  trigger <vsView, view>

```

- Using uniform reliable broadcast instead of best effort broadcast in the previous algorithms does not ensure the uniformity of the message delivery