

# Distributed algorithms

Pierre Colson

Wednesday 05 January 2022

## Contents

<b>General</b>	<b>2</b>
Fair-loss links . . . . .	2
Stubborn links . . . . .	2
Reliable (Perfect) links . . . . .	2
<b>Reliable Broadcast</b>	<b>3</b>
Best-effort Broadcast (beb) . . . . .	3
Reliable Broadcast (rb) . . . . .	3
Uniform Reliable Broadcast (urb) . . . . .	4
<b>Causal Broadcast</b>	<b>4</b>
Causality . . . . .	4
Causal broadcast . . . . .	4
Reliable Causal Broadcast (rcb) . . . . .	4
Uniform Causal Broadcast (ucb) . . . . .	4
Reliable Causal Order Broadcast (rco) . . . . .	5
<b>Total Order Broadcast (tob)</b>	<b>5</b>
(Uniform) Consensus . . . . .	6
Total Order (to) . . . . .	6
<b>Shared Memory</b>	<b>6</b>
Regular register . . . . .	6
Atomic Register . . . . .	8
<b>Terminating Reliable Broadcast (TRB)</b>	<b>9</b>

---

Markdown version on [github](#)  
Compiled using [pandoc](#)  
More fiches [here](#)

## General

- The distributed system is made of a finite set of **processes** : each process models a **sequential** program
- Every pair of processes is connected by a **link** through which the processes exchange **messages**
- **Safety** is a property which states that nothing bad should happen
- **Liveness** is a property which states that something good should happen
- Two kinds of failures are mainly considered
  - **Omissions** : The process omits to send messages it is supposed to send
  - **Arbitrary** : The process sends messages it is not supposed to send
- A **correct** process is a process that does not fail (that does not crash)
- A **Failure detector** is a distributed oracle that provides processes with suspicions about crashed processes
  - It is implemented using *timing assumptions*
  - **Perfect** :
    - \* *Strong Completeness* : Eventually, every process that crashes is permanently suspected by every other correct process
    - \* *Strong Accuracy* : No process is suspected before it crashes
  - **Eventually Perfect** :
    - \* *Strong Completeness*
    - \* *Eventually Strong Accuracy* : Eventually, no correct process is ever suspected

## Fair-loss links

- **FL1. Fair-loss** : If a message is sent infinitely often by  $p_i$  to  $p_j$  and neither  $p_i$  or  $p_j$  crashes then  $m$  is delivered infinitely often by  $p_j$
- **FL2. Finite duplication** : If a message  $m$  is sent a finite number of times by  $p_i$  to  $p_j$ ,  $m$  is delivered a finite number of times by  $p_j$
- **FL3. No creation** : No message is delivered unless it was sent

## Stubborn links

- **SL1. Stubborn delivery** : If a process  $p_i$  sends a message  $m$  to a correct process  $p_j$ , and  $p_i$  does not crash, then  $p_j$  delivers  $m$  an infinite number of times
- **SL2. No creation** : No message is delivered unless it was sent

```
Implements: StubbornLinks (sp2p)
Uses : FairLossLinks (flp2p)
upon event <sp2pSend, dest, m> do
  while (true) do
    trigger <flp2pSend, dest, m>
upon event <flp2pDeliver, src, m> do
  trigger <sp2pDeliver, src, m>
```

## Reliable (Perfect) links

- **PL1. Validity** : If  $p_i$  and  $p_j$  are correct
- **PL2. No duplication** : No message is delivered (to a process) more than once
- **PL3. No creation** : No message is delivered unless it was sent
- Roughly speaking, reliable links ensure that messages exchanged between correct processes are *not lost*

```
Implements: PerfectLinks (pp2p)
Uses: StubbornLinks (sp2p)
upon event <Init> do delivered := emptySet
upon event <pp2pSend, dest, m> do
  trigger <sp2pSend, dest, m>
upon event <sp2pDeliver, src, m> do
```

```

if m not in delivered then
  trigger <pp2pDeliver, src, m>
  add m to delivered

```

## Reliable Broadcast

### Best-effort Broadcast (beb)

- **BEB1. Validity** : If  $p_i$  and  $p_j$  are correct then every message broadcast by  $p_i$  is eventually delivered by  $p_j$
- **BEB2. No duplication** : No message is delivered more than once
- **BEB3. No creation** : No messages is delivered unless it was broadcast

**Implements:** BestEffortBroadcast (beb)

**Uses:** PerfectLinks (pp2p)

```

upon event <bebBroadcast, m> do
  forall pi in S do
    trigger <pp2pSend, pi, m>
  upon event <pp2pDeliver, pi, m> do
    trigger <dedDeliver, pi, m>

```

### Reliable Broadcast (rb)

- **RB1** = BEB1
- **RB2** = BEB2
- **RB3** = BEB3
- **RB4. Agreement** : For any message  $m$ , if any correct process delivers  $m$ , then every correct process delivers  $m$

**Implements:** ReliableBroadcast (rb)

**Uses:**

BestEffortBroadcast (beb)  
PerfectFailureDetector (P)

```

upon event <Init> do
  delivered := emptySet
  correct := S
  forall pi in S do from[pi] := emptySet
upon event <rbBroadcast, m> do
  delivered := delivered U {m}
  trigger <rbDeliver, self, m>
  trigger <bebBroadcast, [data, self, m]>
upon event <crash, pi> do
  correct := correct \ {pi}
  forall [pj, m] in from[pi] do
    trigger <bebBroadcast, [data, pj, m]>
upon event <bebDeliver, pi, [data, pj, m]> do
  if m not in delivered then
    delivered := delivered U {m}
    trigger <rbDeliver, pj, m>
    if pi not in correct then
      trigger <bebBroadcast, [data, pj, m]>
  else
    from[pi] := from[pi] U {[pj, m]}

```

## Uniform Reliable Broadcast (urb)

- URB1 = BEB1
- URB2 = BEB2
- URB3 = BEB3
- URB4. **Uniform Agreement** : For any message  $m$ , if any process delivers  $m$ , then every process delivers  $m$

**Implements:** UniformBroadcast (urb)

**Uses:**

BestEffortBroadcast (beb)  
PerfectFailureDetector (P)

```
upon event <Init> do
  correct := S
  delivered := forward := emptySet
  ack[Message] := emptySet
upon event <urbBroadcast, m> do
  forward := forward U {[self, m]}
  trigger <bebBroadcast, [data, self, m]>
upon event <bebDeliver, pi, [data, pj, m]> do
  ack[m] := ack[m] U {pi}
  if [pi, m] not in forward then
    forward := forward U {[pj, m]}
    trigger <bebBroadcast, [data, pj, m]>
upon event (for any [pj, m] in forward) <correct in ack[m]> and <m not in delivered> do
  delivered := delivered U {m}
  trigger <urbDeliver, pj, m>
```

## Causal Broadcast

- A **non-blocking** algorithm using the past
- A **blocking** algorithm using **vector clocks**

### Causality

- Let  $m_1$  and  $m_2$  be any two messages :  $m_1 \rightarrow m_2$  ( $m_1$  causally precedes  $m_2$ ) iff
  - **C1. FIFO order** : Some process  $p_i$  broadcast  $m_1$  before broadcasting  $m_2$
  - **C2. Local order** : Some process  $p_i$  delivers  $m_1$  and then broadcast  $m_2$
  - **C3. Transitivity** : There is a message  $m_3$  such that  $m_1 \rightarrow m_3$  and  $m_3 \rightarrow m_2$

### Causal broadcast

- **CO** : If any process  $p_i$  delivers a message  $m_2$ , then  $p_i$  must have delivered every message  $m_1$  such that  $m_1 \rightarrow m_2$

## Reliable Causal Broadcast (rcb)

- RB1, RB2, RB3, RB4
- CO

## Uniform Causal Broadcast (ucb)

- URB1, URB2, URB3, URB4
- CO

## Reliable Causal Order Broadcast (rco)

**Implements:** ReliableCausalOrderBroadcast (rco)

**Uses :** ReliableBroadcast (rb)

```
upon event <Init> do
  delivered := past := emptySet
upon event <rcoBroadcast, m> do
  trigger <rbBroadcast, [data, past, m]>
  past := past U {[self, m]}
upon event <rbDeliver, pi [data, pastm, m]> do
  if m not in delivered then
    forall [sn, n] in pastm do
      if n not in delivered then
        trigger <rcoDeliver, sn, n>
        delivered := delivered U {n}
        past := past U {[self, n]}
      trigger <rcoDeliver, pi, m>
    delivered := delivered U {m}
    past := past U {[pi, m]}
```

**Implements** ReliableCausalOrderBroadcast (rco)

**Uses:** ReliableBroadcast (rb)

```
upon event <Init> do
  forall pi in S: VC[pi] := 0
  pending := emptySet
upon event <rcoBroadcast, m> do
  trigger <rcoDeliver, self, m>
  trigger <rbBroadcast, [data, VC, m]>
  VC[self] := VC[self] + 1
upon event <rbDeliver, pj, [data, VCm, m]> do
  if pj not self then
    pending := pending U (pj, [data, VCm, m])
  deliver-pending
procedure deliver-pending is
  while (s, [data, VCm, m]) in pending do
    if forall pk: (VC[pk] >= VCm[pk]) do
      pending := pending - (s, [data, VCm, m])
      trigger <rcoDeliver, self, m>
      VC[s] := VC[s] + 1
```

- These algo ensure causal reliable broadcast
- If we replace reliable broadcast with uniform reliable broadcast, these algo would ensure uniform causal broadcast

## Total Order Broadcast (tob)

- In **reliable** broadcast, the processes are free to deliver messages in any order they wish
- In **causal** broadcast, the processes need to deliver messages according to some order (causal order)
  - The order imposed by causal broadcast is however partial : some messages might be delivered in different order by the processes
- In **total order** broadcast, the processes must deliver all messages according to the same order (i.e. the order is now total)
  - This order does not need to respect causality (or event FIFO ordering)
- **RB1. Validity** : If  $p_i$  and  $p_j$  are correct, then every message broadcast by  $p_i$  is eventually delivered by  $p_j$

- **RB2. No duplication** : No message is delivered more than once
- **RB3. No creation** : No message is delivered unless it was broadcast
- **RB4. (Uniform) Agreement** : For any message  $m$ . If a correct (any) process delivers  $m$ , then every correct process delivers  $m$
- **(Uniform) Total order** : Let  $m$  and  $m'$  be any two messages. Let  $p_i$  be any (correct) process that delivers  $m$  without having delivered  $m'$ . Then no (correct) process delivers  $m'$  before  $m$

## (Uniform) Consensus

- In the (uniform) consensus problem the processes propose values and need to agree on one among these values
- **C1. Validity** : Any value decided is a value proposed
- **C2. (Uniform) Agreement** : No two correct (any) processes decide differently
- **C3. Termination** : Every correct process eventually decides
- **C4. Integrity** : Every process decides at most once

## Total Order (to)

**Implements:** TotalOrder (to)

**Uses:**

```

ReliableBroadcast (rb)
Consensus (cons)
upon event <Init> do
  unordered := delivered := emptySet
  wait := false;
  sn := 1
upon event <toBroadcast, m> do
  trigger <rbBroadcast, m>
upon event <rbDeliver, sm, m> and (m not in delivered) do
  unordered := unordered U {(sm, m)}
upon event (unordered not emptySet) and not wait do
  wait := true
  trigger <Propose, unordered>sn
upon event <Decide, decided>sn do
  unordered := unordered \ decided
  ordered := deterministicSort(decided)
  forall (sm, m) in ordered do
    trigger <toDeliver, sm, m>
    delivered := delivered U {m}
  sn := sn + 1
  wait = false

```

## Shared Memory

### Regular register

- Assumes only one writer
- Provides *strong* guarantees when there is no concurrent operations
- When some operations are concurrent, the register provides *minimal* guarantees
- **Read()** returns :
  - *The last value* written if there is no concurrent or failed operations

- Otherwise the last value written on *any* value concurrently written i.e. the input parameter of some `Write()`
- We assume **fail-stop** model
  - Process can fail by crashing (no recovery)
  - Channels are reliable
  - Failure detection is perfect
- We implement a **regular** register
  - Every process  $p_i$  has a local copy of the register value  $v_i$
  - Every process reads **locally**
  - The writer writes **globally**

```

Write(v) at pi
  send [W, w] to all
  forall pj, wait until either
    receive [ack] or
    detect [pj]
  return ok

```

```

Read() at pi
  return vi

```

```

At pi
  when receive [W, w] from pj
    vi := w
    send [ack] to pj

```

- We assume while failure detection is not perfect
  - $P_1$  is the writer and any process can be reader
  - A majority of the process is correct
  - Channels are reliable
- We implement a **regular** register
  - Every process  $p_i$  maintains a local copy of the register  $v_i$ , as well as a sequence number  $sn_i$  and a read timestamp  $rs_i$
  - Process  $p_1$  maintains in addition a timestamp  $ts_1$

```

Write(v) at p1
  ts1 ++
  send [W, ts1, v] to all
  when receive [W, ts1, ack] from majority
    return ok

```

```

Read() at pi
  rsi ++
  send [R, rsi] to all
  when receive [R, rsi, snj, vj] from majority
    v := vj with the largest snj
  return v

```

```

At pi
  when receive [W, ts1, v] from p1
    if ts1 > sni then
      vi = v
      sni := ts1
      send[W, ts1, ack] to p1
  when receive [R, rsj] from pj
    send [R, rsj, sni, vi] to pj

```

## Atomic Register

- An **Atomic Register** provides strong guarantees even when there is concurrency and failures : the execution is equivalent to a sequential and failure-free execution
- Every failed (write) operation appears to be either complete or not to have been invoked at all
- Every complete operation appears to be executed at some instant between its invocation and reply time events
- We implement a **fail-stop 1-N atomic register**
  - Every process maintains a local value of the register as well as a sequence number
  - The writer,  $p_1$ , maintains, in addition a timestamp  $ts_1$
  - Any process can read in the register

```
Write(v) at p1
  ts1++
  send [W, ts1, v] to all
  forall pi wait until either
    receive [ack] or
    detect [pi]
  return ok
```

```
Read() at pi
  send [W, sni, vi] to all
  forall pi wait until either
    receive [ack] or
    suspect [pj]
  return vi
```

```
At pi
  When pi receive [W, ts, v] from pj
    if ts > sni then
      vi := v
      sni := ts
    send [ack] to pj
```

- We implement a **fail-stop N-N atomic register**

```
Write(v) at pi
  send [W] to all
  forall pj wait until either
    receive [W, snj] or
    suspect [pj]
  (sn, id) := (highest snj + 1, i)
  send [W, (sni, id), v] to all
  forall pj wait until either
    receive [W, (sn, id), ack] or
    detect [pj]
  return ok
```

```
Read() at pi
  send [R] to all
  forall pj wait until either
    receive [R, (snj, idj), vj] or
    suspect pj
  v = vj with the highest (snj, idj)
  (sn, id) = highest (snj, idj)
  send [W, (sn, id), v] to all
```



```

forall pj wait until either
  receive [W, (sn, id), ack] or
  detect [pj]
return v

```

At  $p_i$

T1 :

```

when receive [W] from pj
  send [W, sn] to pj
when receive [R] from pj
  send [R, (sn, id), vi] to pj

```

T2 :

```

when receive [W, (snj, idj), v] from pj
  if (snj, idj) > (sn, id) then
    vi := v
    (sn, id) = (snj, idj)
  send [W, (sn, id), ack] to pj
when receive [W, (snj, idj), v] from pj
  if (snj, idj) > (sn, id) then
    vi := v
    (sn, id) := (snj, idj)
  send [W, (sn, id), ack] to pj

```

- From fail-stop to **fail-silent**
  - We assume a majority of correct processes
  - In the 1-N algorithm, the writer writes in a majority using a timestamp determined locally and the reader selects a value from a majority and then imposes this value on a majority
  - In the N-N algorithm, the writers determines first the timestamp using a majority

## Terminating Reliable Broadcast (TRB)

- Like reliable broadcast, terminating reliable broadcast (TRB) is a communication primitive used to disseminate a message among a set of processes in a reliable way
- TRB is however strictly stronger than (uniform) reliable broadcast
- Like with reliable broadcast, correct processes in TRB agree on the set of messages they deliver
- Like with (uniform) reliable broadcast, every correct process in TRB delivers every message delivered by any correct process
- Unlike with reliable broadcast, every correct process delivers a message, even if the broadcaster crashes
- The problem is defined for a specific broadcaster process  $p_i = src$  (known by all processes)
  - Process  $src$  is supposed to broadcast a message  $m$  (distinct from  $\varphi$ )
  - The other processes need to deliver  $m$  if  $src$  is correct but may deliver  $\varphi$  if  $src$  crashes
- **TRB1. Integrity** : If a process delivers a message  $m$ , then either  $m$  is  $\varphi$  or  $m$  was broadcasted by  $src$
- **TRB2. Validity** : If the sender  $src$  is correct and broadcasts a message  $m$ , then  $src$  eventually delivers  $m$
- **TRB3. (Unifrom) Agreement** : For any message  $m$ , if a correct (any) process delivers  $m$ , then every correct process delivers  $m$
- **TRB4. Termination** : Every correct process eventually delivers exactly one message

**Implements:** trbBroadcast (trb)

**Uses:**

```

BestEffortBroadcast (beb)
PerfectFailureDetector (P)
Consensus (cons)

```

upon event <Init> do

```

prop := 0
correct := S
upon event <trbBroadcast, m> do
  trigger <bebBroadcast, m>
upon event <crash, src> and (prop = 0) do
  prop := phi
upon event <bebDeliver, src, m> and (prop = 0) do
  prop := m
upon event (prop != 0) do
  trigger <Propose, prop>
upon event <Decide, decision> do
  trigger <trbDeliver, src, decision>

```

- We give an algorithm that implements  $P$  unsigned TRB. More precisely, we assume that every process  $p_i$  can use an infinite number of instances of TRB where  $p_i$  is the sender  $src$ 
  1. Every process  $p_i$  keeps on trbBroadcasting messages  $m_{i1}, m_{i2}$  etc
  2. If a process  $p_k$  delivers  $\varphi_i$ ,  $p_k$  suspects  $p_i$