Distributed algorithms

Pierre Colson

Wednesday 05 January 2022

Contents

Markdown verion on *github*
Compiled using *pandoc* and `gpdf` *script*
More fiches *here*

General

- The distributed system is made of a finite set of **processes** : each process models a **sequencial** program
- Every pair of processes is connected by a **link** through which the processes exchange **messages**
- **Safety** is a property which states that nothing bad should happen
- **Liveness** is a property which states that something good should happen
- Twos kinds of failures are mainly considered
  - **Omissions** : The process omits to send messages it is supposed to send
  - **Arbitrary** : The process sends messages it is not supposed to send
- A **correct** process is a process that does not fail (that does not crash)

Failure Detector

- **Indication** : <Crash ,p>
- A **Failure detector** is a distributed oracle that provides processes with suspicions about crashed processes
  - It is implemented using *timing assumptions*
- **Perfect Failure Detector** :
  - **Strong Completness** : Eventually, every process that crashes is permanantly suspected by every other correct process
  - **Strong Accuracy** : No process is suspected before it crashes
- **Eventually Perfect Failure Dectector** :
  - **Strong Completness**
  - **Eventually Strong Accuracy** : Eventually, no correct process is ever suspected

Fair-loss Links

- **FL1. Fai-loss** : If a message is sent infinitely often by p_i to p_jn and neither p_i or p_j crashes then m is delivered infinitely often by p_j
- **FL2. Finite duplication** : If a message m is sent a finite number if times by p_i to p_j, m is delivered a finite number of times by p_j
- **FL3. No creation** : No message is delivered unless it was sent

Stubborn Links

- **SL1. Stubborn delivery** : If a process p_i sends a message m to a correct process p_j, and p_i does not crash, then p_j delivers m an infinite number of times
- **SL2. No creation** : No message is delivered unless it was sent

```
Implements: StubbornLinks (sp2p)
Uses : FairLossLinks (flp2p)
upon event <sp2pSend, dest, m> do
  while (true) do
    trigger <flp2pSend, dest, m>
upon event <flp2pDeliver, src, m> do
  trigger <sp2pDeliver, src, m>
```

Reliable (Perfect) Links

- **PL1. Validity** : If p_i and p_j are correct
- **PL2. No duplication** : No message is delivered (to a process) more than once
- **PL3. No creation** : No message is delivered unless it was sent
- Roughly speaking, reliable links ensure that messages exchanged between correct processes are *not lost*

```
Implements: PerfectLinks (pp2p)
Uses: StubbornLinks (sp2p)
```

```
upon event <Init> do delivered = emptySet
upon event <pp2pSend, dest, m> do
  trigger <sp2Send, dest, m>
upon event <sp2pDeliver, src, m> do
  if m not in delivered then
    trigger <pp2pDeliver, src, m>
    add m to delivered
```

Reliable Broadcast

Best-effort Broadcast (beb)

- **Request** : <bebBroadcast, m>
- **Indication** : <bebDeliver, src, m>
- **BEB1. Validity** : If p_i and p_j are correct then every message broadcast by p_i is eventually delivered by p_j
- **BEB2. No duplication** : No message is delivered more than once
- **BEB3. No creation** : No messages is delivered unless it was broadcast

```
Implements: BestEffortBroadcast (beb)
Uses: PerfectLinks (pp2p)
upon event <bebBroadcast, m> do
  forall pi in S do
    trigger <pp2pSend, pi, m>
upon event <pp2pDeliver, pi, m> do
  trigger <dedDeliver, pi, m>
```

Reliable Broadcast (rb)

- **Request** : <rbBroadcast, m>
- **Indication** : <rbDeliver, src, m>
- **RB1** = BEB1
- **RB2** = BEB2
- **RB3** = BEB3
- **RB4. Agreement** : For any message m, if any correct process delivers m, then every correct process delivers m

```
Implements: ReliableBroadcast (rb)
Uses:
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)
upon event <Init> do
  delivered = emptySet
  correct = S
  forall pi in S do from[pi] = emptySet
upon event <rbBroadcast, m> do
  delivered = delivered U {m}
  trigger <rbDeliver, self, m>
  trigger <bebBroadcast, [data, self, m]>
upon event <crash, pi> do
  correct = correct \{pi}
  forall [pj, m] in from[pi] do
    trigger <bebBroadcast, [data, pj, m]>
upon event <bebDeliver, pi, [data, pj, m]> do
  if m not in delivered then
```

```
    delivered = delivered U {m}
    trigger <rbDeliver, pj, m>
    if pi not in correct then
      trigger <bebBroadcast, [data, pj, m]>
    else
      from[pi] = from[pi] U {[pj, m]}
```

Uniform Reliable Broadcast (urb)

- **Request** : <urbBroadcast, m>
- **Indication** : <urbDeliver, src, m>
- **URB1** = BEB1
- **URB2** = BEB2
- **URB3** = BEB3
- **URB4. Uniform Agreement** : For any message m, if any process delivers m, then every correct process delivers m

```
Implements: UniformBroadcast (urb)
Uses:
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)
upon event <Init> do
  correct = S
  delivered = forward = emptySet
  ack[Message] = emptySet
upon event <urbBroadcast, m> do
  forward = forward U {[self, m]}
  trigger <bebBroadcast, [data, self, m]>
upon event <bebDeliver, pi, [data, pj, m]> do
  ack[m] = ack[m] U {pi}
  if [pi, m] not in forward then
    forward = forward U {[pj, m]}
    trigger <bebBroadcast, [data, pj, m]>
upon event (for any [pj, m] in forward) <correct in ack[m]> and <m not in delivered> do
  delivered = delivered U {m}
  trigger <urbDeliver, pj, m>
```

Causal Broadcast

- A **non-blocking** algorithm using the past
- A **blocking** algorithm using **vector clocks**

Causality

- Let m_1 and m_2 be any two messages : $m_1 \rightarrow m_2$ (m_1 causally precedes m_2) iff
  - **C1. Fifo order** : Some process p_i broadcast m_1 before broadcasting m_2
  - **C2. Local order** : Some process p_i delivers m_1 and then broadcast m_2
  - **C3. Transitivity** : There is a message m_3 such that $m_1 \rightarrow m_3$ and $m_3 \rightarrow m_2$

Causal Broadcast

- **Request** : <coBroadcast, m>
- **Indication** : <coDeliver, src, m>
- **CO** : If any process p_i delivers a message m_2, then p_i must have delivered every message m_1 such that $m_1 \rightarrow m_2$

Reliable Causal Broadcast (rcb)

- **Request** : <rcoBroadcast, m>
- **Indication** : <rcoDeliver, src, m>
- **RB1**, **RB2**, **RB3**, **RB4**
- **CO**

Uniform Causal Broadcast (ucb)

- **Request** : <ucoBroadcast, m>
- **Indication** : <ucoDeliver, src, m>
- **URB1**, **URB2**, **URB3**, **URB4**
- **CO**

Reliable Causal Order Broadcast (rco)

```
Implements: ReliableCausalOrderBroadcast (rco)
Uses : ReliableBroadcast (rb)
upon event <Init> do
  delivered = past = emptySet
upon event <rcoBroadcast, m> do
  trigger <rbBroadcast, [data, past, m]>
  past = past U {[self, m]}
upon event <rbDeliver, pi [data, pastm, m]> do
  if m not in delivered then
    forall [sn, n] in pastm do
      if n not in delivered then
        trigger <rcoDeliver, sn, n>
        delivered = delivered U {n}
        past = past U {[self, n]}
    trigger <rcoDeliver, pi, m>
    delivered = delivered U {m}
    past = past U {[pi, m]}

Implements ReliableCausalOrderBroadcast (rco)
Uses: ReliableBroadcast (rb)
upon event <Init> do
  forall pi in S: VC[pi] = 0
  pending = emptySet
upon event <rcoBroadcast, m> do
  trigger <rcoDeliver, self, m>
  trigger <rbBroadcast, [data, VC, m]>
  VC[self] = VC[self] + 1
upon event <rbDeliver, pj, [data, VCm, m]> do
  if pj not self then
    pending = pending U (pj, [data, VCm, m])
    deliver-pending
procedure deliver-pending is
  while (s, [data, VCm, m]) in pending do
    if forall pk: (VC[pk] >= VCm[pk]) do
      pending = pending - (s, [data, VCm, m])
      trigger <rcoDeliver, self, m>
      VC[s] = VC[s] + 1
```

- These algo ensure causal reliable broadcast

- If we replace reliabe broadcast with uniform reliabe broadcast, these algo would ensure uniform causal broadcast

Total Order Broadcast (to)

- In **reliable** broadcast, the processes are free to deliver messages in any order they wish
- In **causal** broadcast, the processes need to deliver messages according to some order (causal order)
  - The order imposed by causal broadcast is however partial : some messages might be delivered in different order by the processes
- In **total order** broadcast, the processes must deliver all messages according to the same order(i.e. the order is now total)
  - This order does not need to respect causality (or event FIFO ordering)
- **Request** : <toBroadcast, m>
- **Indication** : <toDeliver, src, m>
- **RB1. Validity** : If p_i and p_j are correct, then every message broadcast by p_i is eventually delivered by p_j
- **RB2. No duplication** : No message is delivered more than once
- **RB3. No creation** : No message os delivered unless it was broadcast
- **RB4. (Uniform) Agreement** : For any message m. If a correct (any) process delivers m, then every correct process delivers m
- **(Uniform) Total order** : Let m and m' be any two messages. Let p_i be any (correct) process that delivers m without having delivered m'. Then no (correct) process delivers m' before m
- Uses consensus (see next chapter)

```
Implements: TotalOrder (to)
Uses:
  ReliableBroadcast (rb)
  Consensus (cons)
upon event <Init> do
  unordered = delivered = emptySet
  wait = false;
  sn = 1
upon event <toBroadcast, m> do
  trigger <rbBroadcast, m>
upon event <rbDeliver, sm, m> and (m not in delivered) do
  unordered = unordered U {(sm, m)}
upon event (unordered not emptySet) and not wait do
  wait = true
  trigger <Propose, unordered>sn
upon event <Decide, decided>sn do
  unordered = unordered \ decided
  ordered = deterministicSort(decided)
  forall (sm, m) in ordered do
    trigger <toDeliver, sm, m>
    delivered = delivered U {m}
  sn = sn + 1
  wait = false
```

Consensus

- In the (uniform) consensus problem the processes propose values and need to agree on one among these values
- **Request** : <Propose, v>
- **Indication** : <Decide, v'>

- **C1. Validity** : Any value decided is a value proposed
- **C2. (Uniform) Agreement** : No two correct (any) processes decide differently
- **C3. Termination** : Every correct process eventually decides
- **C4. Integrity** : Every process decides at most once, No process decides twice

Algorithm 1 (cons)

- A *P*-based (fail-stop) consensus algorithm
- The processes exchange and update proposals in rounds and decide on the value of the non-suspected process with the smallest id
- The processes go through rounds incrementally (1 to n) : in each round, the process with the id corresponding to that round is the leader of the round
- The leader of a round decides its current proposal and broadcast it to all
- A process that is not leader in a round waits to deliver the proposal of the leader in that round to adopt, or to suspect the leader

```
Implements: Consensus (cons)
Uses:
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)
upon event <Init> do
  suspected = emptySet
  round = 1
  currentProposal = null
  broadcast = delivered[] = false
upon event <crash, pi> do
  suspected := suspected U {pi}
upon event <Propose, v> do
if currentPoposal = null then
  currentPoposal = v
upon event <bebDeliver, p_round, value> do
  currentPoposal = value
  delivered[round] = true
upon event delivered[round] = true or p_round in suspected do
  round = round + 1
upon event p_round = self and broadcasted = false and currentPoposal not null do
  trigger <Decide, currentPoposal>
  trigger <bebBroadcast, currentPoposal>
  broadcast = true
```

Algorithm 2 (ucons)

- A *P*-based (fail-stop) uniform consenus algorithm
- The processes exchange and update proposal in rounds, and after n rounds decide on the current proposal value
- Implements uniform consensus
- The processes go through rounds incrementally (1 to n): in each round I, process p_I sends its current-Proposal to all
- A process adpots any currentProposal it receives
- Processes decide on their currentProposal values at the end of round n

```
Implements: Uniform Consensus (ucons)
Uses :
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)
```

```
upon event <Init> do
  suspected = emptySet
  round = 1
  currentProposal = null
  broadcast = delivered[] = false
  decided = false
upon event <crash, pi> do
  suspected = suspected U {pi}
upon event <Propose, v> do
  if currentProposal = null then
    currentProposal = v
upon event <bebDeliver, p_round, value> do
  currentProposal = value
  delivered[round] = true
upon event delivered[round] = true or p_round in suspected do
  if round = n and decided = false then
    trigger <Decide, currentProposal>
    decided = true
  else
    round = round + 1
upon event p_round = self and broadcast = false and currentProposal not null do
  trigger <bebBroadcast, currentProposal>
  broadcast = true
```

Algorithm 3

- A $<>P$-based uniform algorithm assuming a correct majority
  - $<>P \Rightarrow$ Eventually Perfect Failure Detector
- The processes alternate in the role of a coordinator until one of them succeeds in imposing a decision
- A uniform consensus algorithm assuming
  - A correct majority
  - A $<>P$ failure detector
- The processes alternate in the role of a phase coordinator until one of them suceeds in imposing a decision
- $<>P$ ensures
  - **Strong completeness** : eventually every process that crashes is permanantly suspected by all correct processes
  - **Eventual strong accuracy** : eventually no correct process is suspected by any process
    * Strong accuracy holds only after finite time
  - Correct processes may be falsely suspected a finite number of times
  - This breaks consensus algorithm 1 and 2
- This algorithm is also round based : process move incrementally from one round to the other
- Process p_i is leader in every round k such that $k mod N = i$
- In such a round, p_i tries to decide
  - p_i succeeds if it is not suspected (process that suspect p_i inform p_i and move to the next round, p_i does so as well)
  - If p_i succeeds, p_i uses a reliable broadcast to send the decision to all (the reliability of the broadcast is important here to preclude the case where p_i crashes, some other processes delivers the message and stop while rest keeps going without majority)
  1. p_i selects among a majority the latest adopted value (lastest with respect to the round in which the value is adopted)
  2. p_i imposes that value at a majority : any process in that majority adopts that value - p_i fails if it is suspected
  3. p_i decides and broadcasts the decision to all

Shared Memory

Regular Register

- Assumes only one writer
- Provides *strong* guarantees when there is no concurrent operations
- When some operations are concurrent, the register provides *minimal* guarantees
- `Read()` returns :
  - *The last value* written if there is no concurrent or failed operations
  - Otherwise the last value written on *any* value concurrently written i.e. the input parameter of some `Write()`
- We assume **fail-stop** model
  - Process can fail by crashing (no recovery)
  - Channels are reliable
  - Failure detection is perfect
- We implement a **regular** register
  - Every process p_i has a local copy of the register value v_i
  - Every process reads **locally**
  - The writer writes **globally**

```
Write(v) at pi
  send [W, w] to all
  forall pj, wait until either
    receive [ack] or
    detect [pj]
  return ok

Read() at pi
  return vi

At pi
  when receive [W, w] from pj
    vi = v
    send [ack] to pj
```

- We assume while failure detection is not perfect
  - P_1 is the writer and any process can be reader
  - A mojority of the process is correct
  - Channels are reliable
- We implement a **regular** register
  - Every process p_i maintains a local copy of the register v_i, as well as a sequence number sn_i and a read timestamp rs_i
  - Process p_1 maintains in addition a timestamp ts_1

```
Write(v) at p1
  ts1 ++
  send [W, ts1, v] to all
  when receive [W, ts1, ack] from majority
    return ok

Read() at pi
  rsi ++
  send [R, rsi] to all
  when receive [R, rsi, snj, vj] from majority
    v = vj with the largest snj
    return v

At pi
```

```
  when receive [W, ts1, v] from p1
    if ts1 > sni then
      vi = v
      sni = ts1
      send[W, ts1, ack] to p1
  when receive [R, rsj] from pj
    send [R, rsj, sni, vi] to pj
```

Atomic Register

- An **Atomic Register** provides strong guarantees event when there is concurrency and failures : the execution is equivalent to a sequencial and failure-free execution

- Every failed (write) operation appears to be either complete or not to have been invoked at all

- Every complete operation appears to be executed at some instant between its invocation and reply time events

- We implement a **fail-stop** 1-N **atomic register**

  - Every process maintais a local value of the register as well as a sequence number
  - The writer, p_1, maintains, in addition a timestamp ts_1
  - Any process can read in the register

```
Write(v) at p1
  ts1++
  send [W, ts1, v] to all
  forall pi wait until either
    receive [ack] or
    detect [pi]
  return ok
```

```
Read() at pi
  send [W, sni, vi] to all
  forall pi wait until either
    receive [ack] or
    suspect [pj]
  return vi
```

```
At pi
  When pi receive [W, ts, v] from pj
    if ts > sni then
      vi = v
      sni = ts
    send [ack] to pj
```

- We implement a **fail-stop** N-N **atomic register**

```
Write(v) at pi
  send [W] to all
  forall pj wait until either
    receive [W, snj] or
    suspect [pj]
  (sn, id) = (highest snj + 1, i)
  send [W, (sni, id), v] to all
  forall pj wait until either
    receive [W, (sn, id), ack] or
    detect [pj]
  return ok
```

```
Read() at pi
  send [R] to all
  forall pj wait until either
    recieve [R, (snj, idj), vj] or
    suspect pj
  v = vj with the highest (snj, idj)
  (sn, id) = highest (snj, idj)
  send [W, (sn, id), v] to all
  forall pj wait until either
    receive [W, (sn, id), ack] or
    detect [pj]
  return v

At pi
T1 :
  when receive [W] from pj
    send [W, sn] to pj
  when receive [R] from pj
    send [R, (sn, id), vi] to pj


T2 :
  when receive [W, (snj, idj), v] from pj
    if (snj, idj) > (sn, id) then
      vi = v
      (sn, id) = (snj, idj)
    send [W, (sn, id), ack] to pj
  when receive [W, (snj, idj), v] from pj
    if (snj, idj) > (sn, id) then
      vi = v
      (sn, id) = (snj, idj)
    send [W, (sn, id), ack] to pj
```

- From fail-stop to **fail-silent**
    - We assume a mojority of correct processes
    - In the 1-N algorithm, the writer writes in a majority using a timestamp determined locally and the reader selects a value from a majority and then imposes this value on a majority
    - In the N-N algorithm, the writers determines first the timestamp using a majority

Terminating Reliable Broadcast (trb)

- Like reliable broadcast, terminating reliable broadcast (TRB) is a communication primitive used to disseminate a message among a set of processes in a reliable way
- TRB is however strictly stronger than (uniform) reliable broadcast
- Like with reliable broadcast, correct processes in TRB agree on the set of messages they deliver
- Like with (uniform) reliable broadcast, every correct process in TRB delivers every message delivered by any correct process
- Unlike with reliable broadcast, every correct process delivers a message, event if the broadcaster crashes
- The problem is defined for a specific broadcaster process $p_i = src$ (know by all processes)
    - Process src is supposed to broadcast a message m (distinct from φ)
    - The other processes need to deliver m if src is correct but may deliver φ is src crashes
- **Request** : <trbBroadcast, m>
- **Indication** : <trbDeliver, p, m>
- **TRB1. Integrity** : If a prcess delivers a message m, then either m is φ or m was broadcasted by src
- **TRB2. Validity** : If the sender src is correct and broadcasts a message m, then src eventually delivers m

- **TRB3. (Uniform) Agreement** : For any message m, if a correct (any) process delivers m, then every correct process delivers m
- **TRB4. Termination** : Every correct process eventually delivers exactly one message

```
Implements: trbBroadcast (trb)
Uses:
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)
  Consensus (cons)
upon event <Init> do
  prop = 0
  correct = S
upon event <trbBroadcast, m> do
  trigger <bebBroadcast, m>
upon event <crash, src> and (prop = 0) do
  prop = phi
upon event <bebDeliver, src, m> and (prop = 0) do
  prop = m
upon event (prop not 0) do
  trigger <Propose, prop>
upon event <Decide, decision> do
  trigger <trbDeliver, src, decision>
```

- We give an algorithm that implements P unsign TRB. More precisely, we assume that every process p_i can use an infinite number of instances of TRB where p_i is the sender *src*
  1. Every process p_i keeps on trbBroadcasting messages m_{i1}, m_{i2} etc
  2. If a process p_k delivers φ_i, p_k suspects p_i

Non-Blocking Atomic Commit (nbac)

- A **transaction** is an atomic program describing a sequence of accesses to shared and distributed information
  – Can be determined either by **commiting** or **aborting**
- **Atomicity** : a transaction either performs entirely or none at all
- **Consistency** : a transaction transforms a consistent state into another consitent state
- **Isolation** : a transaction appears to be executed in isolation
- **Durability** : the effects of a transaction that commits are permanent
- As in consensus, every process has an initial value 0 (no) or 1 (yes) and must decide on a final value 0 (abort) or 1 (commit)
- The proposition means the ability to commit the transaction
- The decision reflects the contract with the user
- Unlike consensus, the processes here seek to decide 1 but every process has a veto right
- **Request** : <Propose, v>
- **Indication** : <Decide, v'>
- **NBAC1. Agreement** : No two processes decide differently
- **NBAC2. Termination** : Every correct process eventually decides
- **NBAC3. Commit-Validity** : 1 can only be decided if all process propose 1
- **NBAC4. Abort-Validity** : 0 can only be decided if some process crashes of votes 0

```
Implements: nonBlockingAtomicCommit (nbac)
Uses:
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)
  UniformConsensus (uniCons)
upon event <Init> do
```

```
    prop = 1
    delivered = emptySet
    correct = PI
upon event <crash, pi> do
  correct = correct \ {PI}
upon event <Propose, v> do
  trigger <bebBroadcast, v>
upon event <bebDeliver, pi, v> do
  delivered = delivered U {pi}
  prop = prop * v
upon event correct \ delivered = empty do
  if correct different PI then
    prop = 0
  trigger <uncPropose, prop>
upon event <uncDecide, decision> do
  trigger <Decide, decision>
```

Group Menbership and View Synchronous Communication

Group Menbership (gmp)

- In many distributed applications, processes neet to know which processes are **participating** in the computation and which are not
- Failure detector provide such information; however that information is **not coordinated** event if the failure detector is perfect
- Like with a failure detector, the processes are informed about failures, we say that the precesses **install views**
- Like with a perfect failure detector, the processes have accurate knowledge about failures
- Unlike with a perfect failure detector, the information about failures are **coordinated** : the processes install the same sequence of views
- **Indication** : <membView, V>
- **Memb1. Local Monotonicity** : if a process installs view (j, M) after installing (k, N), then $j > k$ and $M < N$
- **Memb2. Agreement** : no two porcesses install views (j, M) and (j, M') such that $M \neq M'$
- **Memb3. Completness** : if a process p crashes, then there is an integer j such that every correct porcess eventtually installs view (j, M) such that $p \notin M$
- **Memb4. Accuracy** : if some process installs view (i, M) and $p \notin M$ then p has crashed

```
Implements: groupMenberShip (gmp)
Uses:
  PerfectFailureDetector (P)
  UniformConsensus (Ucons)
upon event <Init> do
  view = (O, S)
  correct = S
  wait = false
upon event <crash, pi> do
  correct = corrext \{pi}
upon event (correct < view.memb) and (wait = false) do
  wait = true
  trigger <ucPropose, (view.id + 1, correct)>
upon event <ucDecide, (id, memb)> do
  view = (id, memb)
  wait = false
```

```
trigger <membView, view>
```

View Synchrony (vsc)

- **View synchronous broadcast** is an abstraction that results from the combination of group membership and reliable broadcast
    - Ensures that the delivery of messages is coordinated with the installation of views
- **Request** : <vsBroadcast, m> (<vsBlockOk>)
- **Indication** : <vsDeliver, src, m>, <vsView, V> (<vsBlock>)
- **Memb1**, **Memb2**, **Memb3**, **Memb4**, **RB1**, **RB2**, **RB3**, **RB4**
- **VS** : a message is **vsDelivered** in the view where it is **vsBroadcast**
- If the application keeps **vsBroadcasting** messages, the **view synchrony** abstraction might never be able to **vsInstall** a new view, the abstraction would be impossible
    - We introduce a specific event for the abstraction to block the application from **vsBroadcasting** messages, this only happends when a process crashes

```
Implements: ViewSynchrony (vs)
Uses:
  GroupMembership (gmp)
  TerminationReliableBroadcast (trb)
  BestEffortBroadcast (beb)
upon event <Init> do
  view = (0, S)
  nextView = ()
  pending = delivered = trbDone = emptySet
  flushing = blocked = false
upon event <vsBroadcast, m> and (blocked = false) do
  delivered = delivered U {m}
  trigger <vsDeliver, self, m>
  trigger <bebBroadcast, [data, view.id, m]>
upon event <bebDeliver, src, [data, vid, m]> do
  if (view.id = vid) and (m not in delivered) and (blocked = false) then
    delivered = delivered U {m}
    trigger <vsDeliver, src, m>
upon event <membView, V> do
  addtoTail(pending, V)
upon event (pending not emptySet) and (flushing = false) do
  nextView = removeFromHead(pending)
  flushing = true
  trigger <vsBlock>
upon event <vsBlockOk> do
  blocked = true
  trbDone = emptySet
  trigger <trbBroadcast, self, (view.id, delivered)>
upon event <trbDeliver, p, (vid, del)> do
  trbDone = trbDone U {p}
  forall m in del and m not in delivered do
    delivered = delivered U {m}
    trigger <vsDeliver, src, m>
upon event (trbDone = view.memb) and (blocked = true) do
  view = nextView
  flushing = blocked = false
  delivered = emptySet
  trigger <vsView, view>
```

```
Implements: ViewSynchrony (vs)
Uses:
  UniformConsensus (uc)
  BestEffortBroadcast (beb)
  PerfectFailureDetector (P)
upon event <Init> do
  view = (0, S)
  correct = S
  flushing = blocked = false
  delivered = dset = emptySet
upon event <vsBroadcast, m> and (blocked = false) do
  delivered = delivered U {m}
  trigger <vsDeliver, self, m>
  trigger <bebBroadcast, [data, view.id, m]>
upon event <bebDeliver, src, [data, vid, m]> do
  if (view.id = vid) and (m not in delivered) and (blocked = false) then
    delivered = delivered U {m}
    trigger <vsDeliver, src, m>
upon event <crash, p> do
  correct = correct \{p}
  if flushing = false then
    flushing = true
    trigger <vsBlock>
upon event <vsBlockOk> do
  blocked = true
  trigger <bebBroadcast, [DSET, view.id, delivered]>
upon event <bebDeliver, src, [DSET, vid, del]> do
  dset = dset U (src, del)
  if forall p in correct, (p, mset) in dset then
    trigger <ucPropose, view.id+1, correct, dset>
upon event <ucDecided, id, memb, vsdset> do
  forall (p, mset) in vsdest and p in memb do
    forall (src, m) in mset and m not in delivered do
      delivered = delivered U {m}
      trigger <vsDeliver, src, m>
    viewx = (id, memb)
    flushing = blocked = false
    dset = delivered = emptySet
    trigger <vsView, view>
```

- Using uniform reliable broadcast instead of best effort broadcast in the previous algorithms does not ensure the uniformity of the message delivery